

Lab Assignment-11

Working with Amazon SageMaker

Amazon SageMaker is a fully managed machine learning service. With SageMaker, data scientists and developers can quickly and easily build and train machine learning models, and then directly deploy them into a production-ready hosted environment.

☐ It provides an integrated Jupyter authoring notebook instance for easy access to your data sources for exploration and analysis, so you don't have to manage servers.

☐ It also provides common machine learning algorithms that are optimized to run efficiently against extremely large data in a distributed environment.

☐ With native support for bring-your-own-algorithms and frameworks, SageMaker offers flexible distributed training options that adjust to your specific workflows.

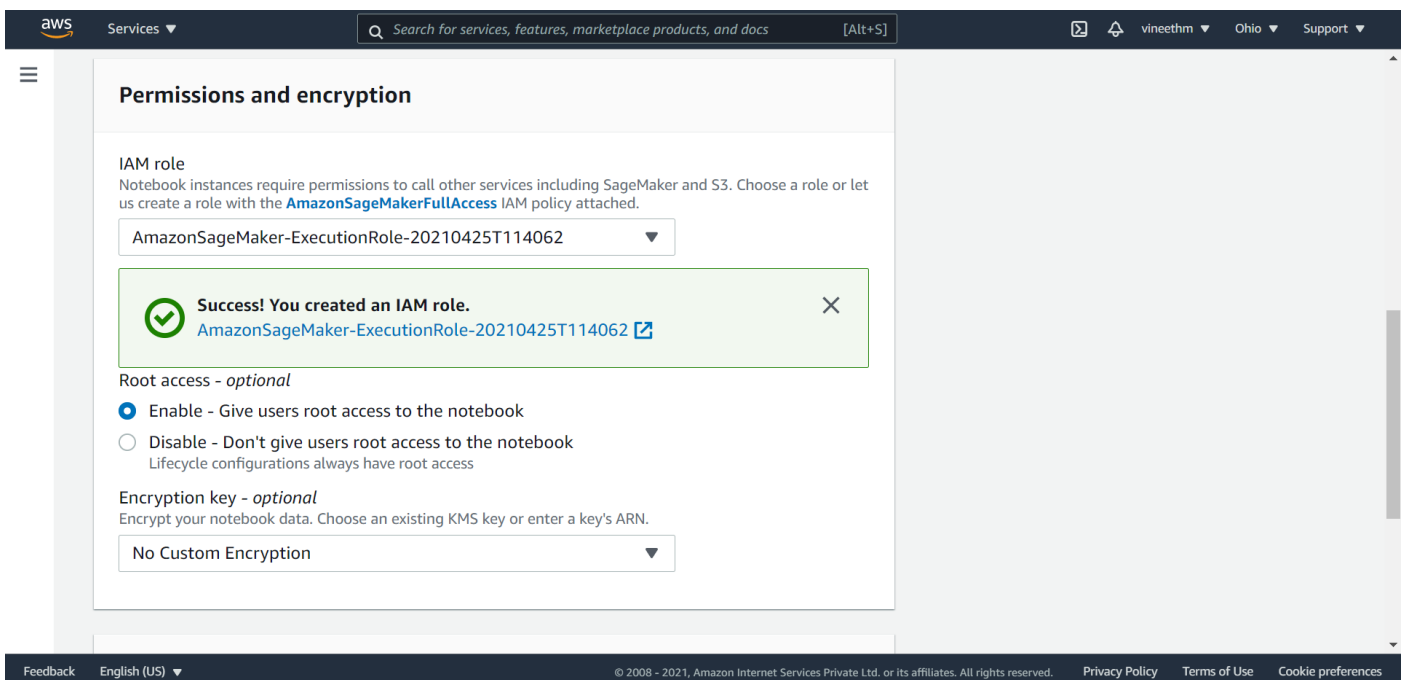
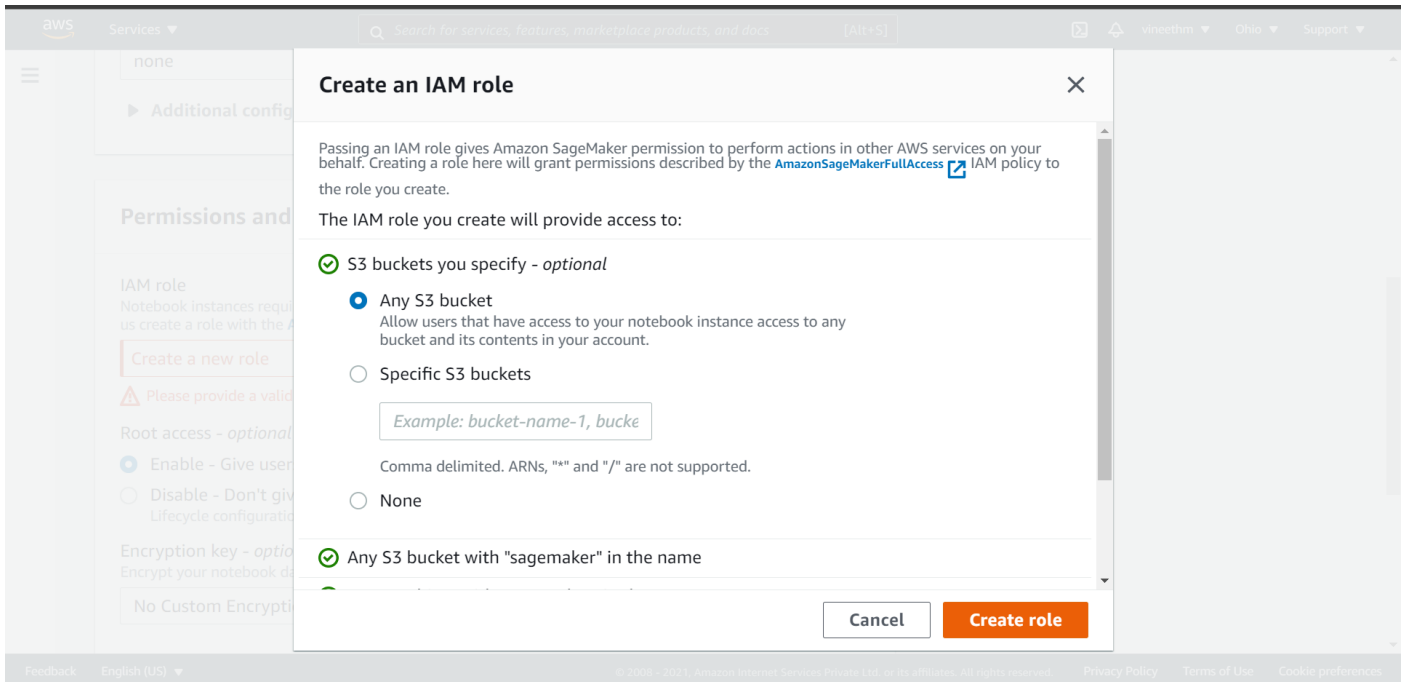
☐ Deploy a model into a secure and scalable environment by launching it with a few clicks from SageMaker Studio or the SageMaker console.

☐ Training and hosting are billed by minutes of usage, with no minimum fees and no upfront Commitments

Step-1: Build, train and deploy a Machine Learning model with Amazon SageMaker.

Task-1: Create a notebook instance.

Note: Create an IAM role to enable the SageMaker access S3 buckets.



The screenshot shows the Amazon SageMaker Studio interface. A green banner at the top states: "Success! Your notebook instance is being created. Open the notebook instance when status is InService and open a template notebook to get started." Below this, the "Notebook instances" section shows a table with one instance:

| Name | Instance | Creation time | Status | Actions |
|---------------|--------------|------------------------|---------|---------|
| xbgBoostModel | ml.t2.medium | Apr 25, 2021 06:11 UTC | Pending | - |

Step-2: Preprocess the data in the Jupyter Notebook provisioned by SageMaker.

Task-1: Choose the kernel environment “conda_python3”.

The screenshot shows the Jupyter Notebook interface with the kernel environment set to "conda_python3". The code in the notebook is as follows:

```
[1]: # import libraries
import boto3, re, sys, math, json, os, sagemaker, urllib.request
from sagemaker import get_execution_role
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import Image
from IPython.display import display
from time import gmtime, strftime
from sagemaker.predictor import csv_serializer

# Define IAM role
role = get_execution_role()
prefix = 'sagemaker/DEMO-xgboost-dm'
containers = {
    'us-west-2': '433757028832.dkr.ecr.us-west-2.amazonaws.com/xgboost:latest',
    'us-east-1': '811284229777.dkr.ecr.us-east-1.amazonaws.com/xgboost:latest',
    'us-east-2': '825641698319.dkr.ecr.us-east-2.amazonaws.com/xgboost:latest',
    'eu-west-1': '685385470294.dkr.ecr.eu-west-1.amazonaws.com/xgboost:latest'
} # each region has its XGBoost container

my_region = boto3.session.Session().region_name # set the region of the instance
print('Success - the MySageMakerInstance is in the ' + my_region + ' region. You will use the ' + containers[my_region] + ' container for your SageMaker endpoint.')

[ ]:
```

```
[6]: bucket_name = 'sagemaker.xgboost.demo' # <--- CHANGE THIS VARIABLE TO A UNIQUE NAME FOR YOUR BUCKET
s3 = boto3.resource('s3')
try:
    if my_region == 'us-east-1':
        s3.create_bucket(Bucket=bucket_name)
    else:
        s3.create_bucket(Bucket=bucket_name, CreateBucketConfiguration={ 'LocationConstraint': my_region })
    print('S3 bucket created successfully')
except Exception as e:
    print('S3 error: ',e)
```

S3 bucket created successfully

```
[7]: try:
      urllib.request.urlretrieve ("https://d1.awsstatic.com/tmt/build-train-deploy-machine-learning-model-sagemaker/bank_clean.27f01
      print('Success: downloaded bank_clean.csv.')
    except Exception as e:
      print('Data load error: ',e)

      try:
        model_data = pd.read_csv('./bank_clean.csv',index_col=0)
        print('Success: Data loaded into dataframe.')
      except Exception as e:
        print('Data load error: ',e)
```

Success: downloaded bank_clean.csv.
Success: Data loaded into dataframe.

Step-3: Train the Machine Learning Model

```
[8]: train_data, test_data = np.split(model_data.sample(frac=1, random_state=1729), [int(0.7 * len(model_data))])
      print(train_data.shape, test_data.shape)

(28831, 61) (12357, 61)

[10]: pd.concat([train_data['y_yes'], train_data.drop(['y_no', 'y_yes'], axis=1)], axis=1).to_csv('train.csv', index=False, header=False)
      boto3.Session().resource('s3').Bucket(bucket_name).Object(os.path.join(prefix, 'train/train.csv')).upload_file('train.csv')
      s3_input_train = sagemaker.inputs.TrainingInput(s3_data='s3://{}/{}/train'.format(bucket_name, prefix), content_type='csv')
```

```
[11]: sess = sagemaker.Session()
      xgb = sagemaker.estimator.Estimator(containers[my_region],role, instance_count=1, instance_type='ml.m4.xlarge',output_path='s3://
      xgb.set_hyperparameters(max_depth=5,eta=0.2,gamma=4,min_child_weight=6,subsample=0.8,silent=0,objective='binary:logistic',num_ro
```

[]:

```
[*]: xgb.fit({'train': s3_input_train})

2021-04-25 06:24:56 Starting - Starting the training job...
2021-04-25 06:25:20 Starting - Launching requested ML instancesProfilerReport-1619331896: InProgress
.....
2021-04-25 06:26:20 Starting - Preparing the instances for training.....
```

[]:

```
[*]: xgb.fit({'train': s3_input_train})
```

2021-04-25 06:24:56 Starting - Starting the training job...

2021-04-25 06:25:20 Starting - Launching requested ML instancesProfilerReport-1619331896: InProgress

.....

2021-04-25 06:26:20 Starting - Preparing the instances for training.....

2021-04-25 06:27:40 Downloading - Downloading input data...

2021-04-25 06:28:24 Training - Training image download completed. Training in progress..Arguments: train

[2021-04-25:06:28:25:INFO] Running standalone xgboost training.

[2021-04-25:06:28:25:INFO] Path /opt/ml/input/data/validation does not exist!

[2021-04-25:06:28:25:INFO] File size need to be processed in the node: 3.38mb. Available memory size in the node: 8418.43mb

[2021-04-25:06:28:25:INFO] Determined delimiter of CSV input is ','

[06:28:25] S3DistributionType set as FullyReplicated

[06:28:25] 28831x59 matrix with 1701029 entries loaded from /opt/ml/input/data/train?format=csv&label_column=0&delimiter=,

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 30 extra nodes, 14 pruned nodes, max_depth=5

[0]#011train-error:0.100482

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 30 extra nodes, 14 pruned nodes, max_depth=5

[1]#011train-error:0.099858

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 28 extra nodes, 22 pruned nodes, max_depth=5

[2]#011train-error:0.099754

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 22 extra nodes, 14 pruned nodes, max_depth=5

[3]#011train-error:0.099095

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 28 extra nodes, 12 pruned nodes, max_depth=5

[4]#011train-error:0.098991

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 32 extra nodes, 14 pruned nodes, max_depth=5

[5]#011train-error:0.099303

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 24 extra nodes, 18 pruned nodes, max_depth=5

[6]#011train-error:0.099684

[06:28:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 30 extra nodes, 22 pruned nodes, max_depth=5

[7]#011train-error:0.09906

[06:28:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 28 extra nodes, 20 pruned nodes, max_depth=5

[8]#011train-error:0.098852

[06:28:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 36 extra nodes, 8 pruned nodes, max_depth=5

Note: The required files are created in the S3 bucket.

The screenshot shows the Amazon S3 console interface. On the left, there's a navigation pane with 'Amazon S3' selected. Below it, there are links for 'Buckets', 'Access Points', 'Object Lambda Access Points', 'Batch Operations', 'Access analyzer for S3', 'Block Public Access settings for this account', and 'Storage Lens'. The main content area shows the 'train/' bucket. At the top right of the bucket view, there's a 'Copy S3 URI' button. Below the bucket name, there are two tabs: 'Objects' (selected) and 'Properties'. The 'Objects' tab displays a list of objects. There is one object named 'train.csv'. The table has columns for 'Name', 'Type', 'Last modified', 'Size', and 'Storage class'. The 'train.csv' object is a CSV file, 3.4 MB in size, and was last modified on April 25, 2021, at 11:53:39 UTC+05:30. The storage class is 'Standard'. Above the table, there are buttons for 'Copy URL', 'Delete', 'Actions', 'Create folder', and 'Upload'. There is also a search bar labeled 'Find objects by prefix'.

Step-4: Deploy Model

```
[13]: xgb_predictor = xgb.deploy(initial_instance_count=1, instance_type='ml.m4.xlarge')
```

```
2021-04-25 06:28:41 Uploading - Uploading generated training model
2021-04-25 06:28:41 Completed - Training job completed
Training seconds: 61
Billable seconds: 61
```

```
[13]: xgb_predictor = xgb.deploy(initial_instance_count=1, instance_type='ml.m4.xlarge')
-----!
```

```
[14]: from sagemaker.serializers import CSVSerializer

test_data_array = test_data.drop(['y_no', 'y_yes'], axis=1).values #Load the data into an array
xgb_predictor.serializer = CSVSerializer() # set the serializer type
predictions = xgb_predictor.predict(test_data_array).decode('utf-8') # predict!
predictions_array = np.fromstring(predictions[1:], sep=',') # and turn the prediction into an array
print(predictions_array.shape)

(12357,)
```

```
[ ]:
```

We can notice the endpoint being created.

The screenshot shows the AWS SageMaker console interface. On the left is a navigation sidebar with categories like Ground Truth, Notebook, Processing, Training, Inference, Edge Manager, Augmented AI, and AWS Marketplace. The 'Inference' section is expanded, showing options like Compilation jobs, Model packages, Models, Endpoint configurations, **Endpoints** (highlighted), and Batch transform jobs. The main content area is titled 'Amazon SageMaker > Endpoints'. It features a 'Create endpoint' button and a table of endpoints. The table has columns for Name, ARN, Creation time, Status, and Last updated. One endpoint is listed with the name 'xgboost-2021-04-25-06-30-01-868', ARN 'arn:aws:sagemaker:us-east-2:277391238495:endpoint/xgboost-2021-04-25-06-30-01-868', creation time 'Apr 25, 2021 06:30 UTC', and status 'Creating'.

| Name | ARN | Creation time | Status | Last updated |
|---------------------------------|---|------------------------|----------|------------------------|
| xgboost-2021-04-25-06-30-01-868 | arn:aws:sagemaker:us-east-2:277391238495:endpoint/xgboost-2021-04-25-06-30-01-868 | Apr 25, 2021 06:30 UTC | Creating | Apr 25, 2021 06:30 UTC |

Step-5: Evaluate Model performance.

```
[15]: cm = pd.crosstab(index=test_data['y_yes'], columns=np.round(predictions_array), rownames=['Observed'], colnames=['Predicted'])
tn = cm.iloc[0,0]; fn = cm.iloc[1,0]; tp = cm.iloc[1,1]; fp = cm.iloc[0,1]; p = (tp+tn)/(tp+tn+fp+fn)*100
print("\n{0:<20}{1:<4.1f}%\n".format("Overall Classification Rate: ", p))
print("{0:<15}{1:<15}{2:>8}".format("Predicted", "No Purchase", "Purchase"))
print("Observed")
print("{0:<15}{1:<2.0f}% ({2:<}){3:>6.0f}% ({4:<})".format("No Purchase", tn/(tn+fn)*100,tn, fp/(tp+fp)*100, fp))
print("{0:<16}{1:<1.0f}% ({2:<}){3:>7.0f}% ({4:<}) \n".format("Purchase", fn/(tn+fn)*100,fn, tp/(tp+fp)*100, tp))
```

Overall Classification Rate: 89.5%

| Predicted | No Purchase | Purchase |
|-------------|-------------|-----------|
| Observed | | |
| No Purchase | 90% (10769) | 37% (167) |
| Purchase | 10% (1133) | 63% (288) |

[]: |

Step-6: Clean up the resources.

The screenshot shows a Jupyter Notebook window titled 'XGBoost.ipynb'. The code in cell [18] uses the boto3 library to delete a bucket named 'bucket_to_delete'. The output of the cell is a large JSON object representing an HTTP response from Amazon S3, including metadata, headers, and a list of deleted objects.

```
[18]: bucket_to_delete = boto3.resource('s3').Bucket(bucket_name)
      bucket_to_delete.objects.all().delete()
```

```
[18]: [{"ResponseMetadata": {"RequestId": "X8269C69BFPC4F0M",
  "HostId": "eaLIVUadgduElV5S9mo035V341skDSMIYIZvU5kuniJdQsMwcB9OwKA7sVupXQ6b8irNLHT7/Ks=",
  "HTTPStatusCode": 200,
  "HTTPHeaders": {"x-amz-id-2": "eaLIVUadgduElV5S9mo035V341skDSMIYIZvU5kuniJdQsMwcB9OwKA7sVupXQ6b8irNLHT7/Ks=",
    "x-amz-request-id": "X8269C69BFPC4F0M",
    "date": "Sun, 25 Apr 2021 06:39:12 GMT",
    "content-type": "application/xml",
    "transfer-encoding": "chunked",
    "server": "AmazonS3",
    "connection": "close"},
  "RetryAttempts": 0},
  "Deleted": [{"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/pr
ofiler-output/profiler-reports/LowGPUUtilization.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/MaxInitializationTime.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/Dataloader.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/OverallSystemUsage.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/profiler-output/system/incremental/2021042506/16193
32080.algo-1.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/StepOutlier.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-report.html"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/output/model.tar.gz"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/OverallFrameworkMetrics.json"},
    {"Key": "sagemaker/DEMO-xgboost-dm/output/xgboost-2021-04-25-06-24-56-728/rule-output/ProfilerReport-1619331896/profiler-outp
ut/profiler-reports/BatchSize.json"}]}
```